

Efficiency: context tokens vs LLM-Judge
fewer tokens, higher quality (upper-left dominates)

