

Efficiency: Latency vs ROUGE-L
(bubble size $\propto$ avg context tokens, $n=204$)

